

Projecting NFL Wide Receiver Contracts

Term Probabilities and Cap-Share Price for the 2026 Veteran Market

Rento Saijo, Atticus Roop, and Manahil Shoaib

May 1, 2026

Abstract. This paper asks whether public data can predict veteran NFL wide receiver contract term, average annual value as a salary-cap share (AAV%), and the variables that make those forecasts credible. Spotrac contract records are enriched with nflverse player, roster, combine, production, snap, Next Gen Stats, team, and schedule data, plus ESPN transaction dates when matched. Because term and annual price are different negotiation choices, we use a multiclass term model and a conditional AAV% model. Tuned XGBoost performs best on signed 2026 validation contracts, with term log loss 0.537 and AAV% RMSE 0.0057. Scoring 163 unsigned 2026 receivers produces a short-term market: every player has a modal one-year deal, with median expected AAV of \$1.11M.

Introduction and Background

An NFL wide receiver is an offensive player whose main job is to catch passes, but wide receiver contracts are not priced on catches alone. A one-year depth contract, a two-year role-player agreement, and a four-year star extension all belong to the same position group while expressing different levels of team commitment, player leverage, age risk, and market timing. Average annual value (AAV) is reported dollars per contract year; AAV% divides AAV by the salary cap at signing so a deal from 2005 and a deal from 2025 can be compared on the same economic scale.

The research questions are: Can veteran wide receiver contract term be predicted from public information available before signing? Can AAV% be predicted with useful accuracy once term is specified? Which player, production, contract-history, team-context, and market variables explain those predictions? Public contract forecasts are often single dollar estimates, while actual negotiation resembles a menu of term-price scenarios. We therefore use two linked model families: a term classifier over one, two, three, four, and five-or-more years, and a conditional price model that estimates AAV% under each possible term scenario.

Figure 1 previews the empirical reason for this architecture. The top row shows that one-year deals are already common in training (61.0%) and become dominant in signed 2026 validation contracts (84.5%), so term must be evaluated with probability scores rather than raw accuracy. The bottom row shows that longer commitments occupy a different AAV% region, which supports pricing AAV% conditional on a supplied term instead of forcing one model to learn both decisions at once.




Figure 1: Term concentration and term-price separation in the labeled data. Bars show each term class as a share of contracts; boxes show the AAV% distribution within each term. The validation market is more one-year concentrated than training, while longer deals occupy higher cap-share ranges, motivating probability-based term evaluation and a conditional AAV% model.

Data, EDA, and Methods

The contract spine comes from Spotrac wide receiver exports by decade ([Spotrac 2026](#)). Cleaning keeps veteran, non-entry wide receiver deals because rookie-scale contracts follow a separate bargaining mechanism. It also standardizes names and teams, parses economics, converts AAV% to cap share, repairs inconsistent end years, removes yearly rows embedded inside prior multi-year contracts, and builds previous-contract features. It then enriches rows with nflverse player references, rosters, combine measurements, regular receiving statistics, snap counts, Next Gen receiving metrics, team offense, and schedules ([nflverse 2026](#)). ESPN transaction descriptions are matched by player, team, and contract clues when possible; otherwise the workflow uses an explicit March 15 offseason proxy. Prior-season predictors are the two most recent completed regular seasons before that reference date, not a simple calendar lag; the most recent season is weighted twice as heavily as the second prior season, trend variables compare the two seasons when both exist, and availability flags preserve whether each feature family was observed.

The final modeling files contain 2,796 training contracts from 2002-2025, 91 signed 2026 validation contracts, and 163 unsigned 2026 free-agent candidates. The validation set is kept separate for model comparison, while the test file intentionally leaves outcomes blank because those players are the actual unsigned market to be scored. The EDA in Figure 1 shaped two choices. First, the term target is imbalanced enough that a model can look accurate by always predicting one year, so model selection must reward calibrated probabilities through log loss and Brier score. Second, AAV% rises with term, but the relationship is not just a dollar ranking of players; it is a negotiation surface where the same receiver can have a low-probability high-term scenario and a high-probability one-year scenario.

Modeling uses `tidymodels` recipes and engines ([tidymodels 2026](#)). Identifiers, audit dates, signing team, and current-contract outcomes are removed. Categorical predictors receive unknown/novel handling, mode imputation, and dummy encoding; numeric predictors receive median imputation, zero-variance filtering, and normalization. The term task drops current AAV and AAV%; the AAV% task drops current AAV but keeps term because production scoring evaluates price under supplied term scenarios. Screening compared null, linear, penalized, tree, random forest, boosted, neural-network, and LightGBM-style candidates where available. Final tuning used grouped five-fold cross-validation by player, which prevents one player's related contracts from appearing in both assessment and fitting folds. The term model minimizes multinomial log loss because the output is a five-class probability vector; the AAV% model minimizes RMSE because the output is a continuous cap-share price. The selected predictive pair is XGBoost for both tasks ([Chen and Guestrin 2016](#)); lasso multinomial regression and ridge regression provide audit models ([Friedman et al. 2010](#)).

Results and Discussion

On the signed 2026 validation set, the tuned XGBoost term model achieved log loss 0.537, Brier score 0.057, and accuracy 82.4%. The null majority classifier is more accurate because 76 of 91 validation contracts are one-year deals, but it is worse on probability scores; that result confirms the EDA rather than contradicting the model. The useful term output is a probability distribution, not a winner-take-all label. The XGBoost AAV% model achieved RMSE 0.0057, MAE 0.0026, and $R^2 = 0.903$. At the implied 2026 cap, the RMSE is roughly \$1.71M, which is precise enough for market screening while still leaving room for information public data cannot see.

The tuned comparisons in Appendix B also show why the final system keeps both predictive and audit models. XGBoost is the best validation choice for term by log loss and for AAV% by RMSE, but the margins over random forest and ridge-style alternatives are small enough that interpretability checks matter. The lasso term model and ridge AAV% model are therefore not fallback forecasts; they are safeguards that expose whether the stronger nonlinear models are leaning on sensible football and market signals. This is especially important because the strongest-looking accuracy benchmark is the least informative one: a model that predicts one year for everyone can exploit the 2026 class mix without producing useful probabilities for the few players who might command multi-year deals.

The interpretation artifacts in Appendix C support the two-model design. Term importance is led by contract number, years of experience, previous-team/matching indicators, prior timing, age, and previous market value; recent production appears, but less dominantly. Compensation importance is led by weighted receiving yards, yards per game, first downs per game, target share, prior cap share, snap role, and the supplied term scenario. In plain language, the model treats term as a commitment and career-stage decision, while AAV% is more

directly a receiving-value and market-price decision. The penalized audit models point in the same direction, which makes the boosted results easier to trust: the strongest predictive signals are football-plausible rather than artifacts of row order, player identifiers, or post-signing outcomes.

Applied to the unsigned 2026 receiver pool, all 163 scored players have a modal one-year contract, the mean one-year probability is 85.8%, and the mean probability of any multi-year outcome is 14.2%. Price remains right-skewed: mean expected AAV is \$1.49M, median expected AAV is \$1.11M, and the maximum is \$11.52M. Deebo Samuel, Jauan Jennings, and Keenan Allen top the expected-AAV ranking, but even their modal forecast remains one year. This is why the final output is a term-price menu: expected AAV weights all term scenarios by their probabilities, while modal AAV reports the price attached to the most likely term. Player-level menus are available in the [companion dashboard](#), and Appendix D shows both the market board and scenario paths.

The conclusion is that the project is useful less as a single prediction machine than as a disciplined way to separate contract questions. The first research question receives a qualified yes: term is predictable enough to rank probabilities, but the validation distribution makes a hard class label too blunt. The second question receives a stronger yes: AAV% is meaningfully predictable when the model is given a term scenario. The third question is answered by the contrast between feature families. Contract history and career stage explain commitment; receiving production, prior price, role, and term explain compensation. That split makes the final app more interpretable because a user can compare one-year, two-year, and longer paths instead of reading one dollar number as certainty.

Limitations and Future Work

The main limitations are measurement, validation size, and unobserved negotiation context. Signing-date matches cover only part of history, player identity matching is probabilistic for fringe players, and Next Gen receiving coverage is thin before recent seasons. The validation set has only 91 signed 2026 contracts and is concentrated in one-year deals, so small metric differences should not be overinterpreted. Public data also omit guarantees, incentives, option years, injury detail, cap room, depth-chart need, agent leverage, and competing offers. That is especially important for AAV: two deals with the same annual value can differ sharply if one is guaranteed and the other is built from non-guaranteed years or incentives.

Future work should add guarantee and incentive targets, automate contract-data refreshes, test calibration and prediction intervals, and move from player-only scoring to team-player matching. A team-aware version would add cap space, receiver depth chart, offensive style, quarterback quality, and recent target vacancy, then score the same receiver under different destinations. The next evaluation should also ask whether a group assigned a 20% chance of multi-year contracts actually signs multi-year deals about one in five times over repeated markets. The present workflow is reproducible enough for that extension: a reader can rebuild the modeling files from the contract exports and public data pulls, verify the time-based split, inspect the baseline and tuned comparisons, and reproduce the final scenario files used by the dashboard.

Appendix

Appendix A: Data Construction and Feature Coverage

Table 1: Data sources and how each source enters the project.

Source	Role in Project
Spotrac WR contract exports	Contract spine: player, team, start/end years, term, value, AAV, AAV%, age, guarantees/bonuses for parsing.
nflreadr player, roster, combine, stats, snap, NGS, team, and schedule data	Public pre-signing player identity, athletic profile, production, usage, receiver-tracking, team offense, and season-calendar context.
ESPN NFL transactions API	Observed or likely signing dates used to reduce temporal leakage in prior-season feature windows.
Derived market context	Rolling veteran WR market anchors from prior local contract records: count, median term, AAV, AAV%, p75 AAV%, and signing age.

The reproducible data workflow begins with three local Spotrac wide receiver contract exports covering the 2000s, 2010s, and 2020s. Table 1 documents the provenance of each data family. The local contract files define the target rows; the external sources add information that would plausibly be known before signing.

Table 2: Modeled responses and statistical rationale.

Response	Definition	Why This Target
term	Reported contract length in years; deals of five or more years are grouped into the 5+ class.	Represents team commitment and player leverage; treated as a probability distribution because one-year deals dominate the validation market.
aavPerc	Average annual value divided by league salary cap at signing, modeled as a decimal share.	Normalizes dollars across eras and gives a price target that can be translated back to 2026 dollars.

Table 2 defines the two modeled responses. This distinction is the reason the final workflow has two model sets: term is a commitment outcome, while AAV% is a price outcome evaluated under a supplied term scenario.

Table 3: Modeling data after cleaning.

Split	Role	Rows	Players	Years
Train	Signed veteran WR contracts used for fitting and tuning	2,796	1,305	2002-2025
Validate	Signed 2026 veteran WR contracts used for held-out evaluation	91	91	2026
Test	Unsigned 2026 WR free agents scored by the final models	163	163	2026

Table 3 gives the final row universe after cleaning. The split is time-based rather than random because the intended use case is forward projection into an unsigned free-agent market.

Table 4: Key cleaning and temporal-logic decisions in R/clean.R.

Step	Implementation Detail
Contract parsing	Row-bind WR contract files, standardize names/teams, parse numeric economics, convert AAV% to cap share, and repair inconsistent end years with $\text{startYear} + \text{term} - 1$.
Embedded-row removal	Drop yearly rows already covered by an earlier multi-year deal for the same player/team before building contract-history variables.
Entry-like filtering	Exclude contracts at rookieSeason or first observed contracts with $\text{ageAtSigning} \leq 23$ to avoid mixing rookie-scale economics with veteran pricing.
Player matching	Join normalized contract names to nflverse player references, then prefer candidates by signing-age gap, rookie-season plausibility, roster backfill, and position context.
Signing-date reference	Use matched ESPN transaction dates when available; otherwise use an explicit March 15 offseason proxy for the contract start year.
Prior-season windows	Use the two most recent fully completed regular seasons before the reference date, not a mechanical $\text{startYear} - 1 / \text{startYear} - 2$ rule.
Missing-data policy	Keep raw missing season fields as NA, collapse weighted values to the observed season when only one year exists, set missing deltas to 0, and retain availability flags.

Table 4 summarizes the cleaning decisions that make the rows reproducible and reduce leakage. The most important details are the removal of entry-like contracts, the removal of embedded rows already covered by prior multi-year deals, and the use of signing-date-aware prior seasons.

Table 5: Training coverage for external feature families.

Feature Family	Training Coverage
Observed ESPN signing date	60.1%
NFL combine	45.4%
Regular receiving stats	65.8%
Snap usage	51.3%
Next Gen receiving	15.8%
Team context	68.8%
Market context	100.0%

Table 5 reports feature coverage in the training file. This also explains why the final app and paper emphasize universally available contract context first, then production, snap, and Next Gen information when present.

Table 6: Contract-term distribution by modeling split.

Split	Term	Contracts	Share
Training contracts	1	1705	61.0%
Training contracts	2	517	18.5%
Training contracts	3	376	13.4%
Training contracts	4	97	3.5%
Training contracts	5+	101	3.6%
Signed 2026 validation	1	76	83.5%
Signed 2026 validation	2	9	9.9%
Signed 2026 validation	3	3	3.3%
Signed 2026 validation	4	3	3.3%

Table 6 gives the numeric version of the term imbalance shown in the main EDA figure. The validation market is substantially more one-year-heavy than the historical training market.

The two-season feature convention creates **Prev1**, **Prev2**, **Weighted**, and **Delta** versions of regular-stat, snap, Next Gen, team, and market metrics. Weighted features use $(\text{Prev2} + 2 * \text{Prev1}) / 3$ when both seasons exist, fall back to the observed season when only one exists, and preserve availability flags. Deltas are set to zero when a trend cannot be observed, which treats missing trend as neutral while keeping true zeroes distinct in raw metrics.

Appendix B: Model Selection

Table 7: Modeling preprocessing and leakage controls.

Component	Modeling Choice
Shared exclusions	Identifiers, audit dates, signing-team label, and current-contract bookkeeping fields are not predictors.
Term-specific exclusions	Current-contract AAV and AAV% are removed to prevent direct target leakage.
AAV%-specific exclusions	Current-contract AAV is removed; term remains because the final scorer prices each supplied term scenario.
Categorical preprocessing	Convert strings to factors, handle unknown and novel levels, mode-impute, and dummy encode.
Numeric preprocessing	Median-impute, remove zero-variance predictors, and normalize.
Resampling and tuning metrics	Use grouped five-fold cross-validation by playerId; optimize multinomial log loss for term and RMSE for AAV%.

Table 7 records the modeling recipe. These choices are shared by the baseline, tuning, final refit, and prediction scripts, so the exported results can be traced back to a single preprocessing contract.

Table 8: Baseline term models ordered by validation log loss.

Model	Log Loss	ROC AUC	Brier	Accuracy
XGBoost	0.560	0.666	0.058	81.3%
Random forest	0.569	0.673	0.059	80.2%
Lasso	0.588	0.645	0.060	81.3%
Elastic net	0.613	0.618	0.061	81.3%
LightGBM	0.652	0.622	0.062	82.4%
Ridge	0.710	0.601	0.065	81.3%
Null majority	0.757	0.500	0.072	83.5%
Multinomial logistic	1.004	0.520	0.075	79.1%
Neural net	1.170	0.546	0.116	72.5%
Decision tree	1.374	0.631	0.069	79.1%

Table 8 shows every successful untuned term baseline. The null classifier is included deliberately because the one-year class imbalance makes it a necessary benchmark: it is accurate but poorly probabilistic.

Table 9: Baseline compensation models ordered by validation RMSE.

Model	RMSE	MAE	R-squared
XGBoost	0.0061	0.0032	0.882
Random forest	0.0062	0.0032	0.876
Ridge	0.0064	0.0040	0.871
LightGBM	0.0066	0.0035	0.862
Linear regression	0.0080	0.0048	0.798
Decision tree	0.0104	0.0041	0.710
Elastic net	0.0108	0.0064	0.741
Neural net	0.0110	0.0070	0.616
Lasso	0.0138	0.0083	0.669
Null mean	0.0177	0.0118	NA

Table 9 shows every successful untuned AAV% baseline. Boosted trees and random forests performed best, while ridge was the strongest linear-style baseline.

Table 10: Tuned model comparison on the signed 2026 validation set.

Task	Model	Validation Metric 1	Validation Metric 2	Validation Metric 3	Cross-Validation Metric
Term	XGBoost	Log loss 0.537	Brier 0.057	Accuracy 82.4%	CV log loss 0.788
Term	Random forest	Log loss 0.545	Brier 0.057	Accuracy 81.3%	CV log loss 0.795
Term	Lasso multinomial	Log loss 0.592	Brier 0.060	Accuracy 81.3%	CV log loss 0.832
AAV%	XGBoost	RMSE 0.0057	MAE 0.0026	R-squared 0.903	CV RMSE 0.0097
AAV%	Random forest	RMSE 0.0061	MAE 0.0031	R-squared 0.880	CV RMSE 0.0105
AAV%	Ridge linear	RMSE 0.0061	MAE 0.0039	R-squared 0.883	CV RMSE 0.0105

Table 10 is the final tuned comparison. The predictive pair is chosen from these validation results; the penalized models remain in the project because they provide interpretable audit trails.

Table 11: Final model families and their roles in the project.

Model	Selected Family	Validation Evidence	Project Role
Predictive term	XGBoost multiclass classifier	Log loss 0.537, Brier 0.057, accuracy 82.4%	Final probability engine for contract term.
Predictive AAV%	XGBoost conditional regressor	RMSE 0.0057, MAE 0.0026, R-squared 0.903	Final scenario-price engine for AAV%.
Interpretable term	Lasso multinomial logistic regression	Log loss 0.592, Brier 0.060, accuracy 81.3%	Audit model with class-specific nonzero coefficients.
Interpretable AAV%	Ridge linear regression	RMSE 0.0061, MAE 0.0039, R-squared 0.883	Audit model with a stable linear coefficient vector.

Table 11 clarifies why there are two sets of final models. The XGBoost pair is the production forecasting pair; the lasso/ridge pair is the transparent audit pair used to check whether the boosted forecasts are relying on football-plausible signals.

Table 12: Selected XGBoost hyperparameters.

Model	mtry	Trees	Minimum N	Tree Depth	Learning Rate	Loss Reduction	Sample Size
Term XGBoost	92	1251	10	4	0.005249	0.00016	0.945583
AAV% XGBoost	57	473	2	3	0.078861	0.00001	0.772218

Table 12 records the selected XGBoost hyperparameters from Bayesian optimization. The tuned term model is deeper and slower-learning; the tuned AAV% model is shallower, which is consistent with a smoother conditional price surface.

Appendix C: Interpretation Artifacts




Figure 2: Top gain-importance features for the predictive term model.

Figure 2 shows the top 25 gain-importance features for the predictive term model. Gain importance is unsigned and non-causal, but it identifies which transformed predictors most improved tree splits. Contract-history, career-stage, and match-quality features dominate, with production appearing after those market-commitment signals.




Figure 3: Top gain-importance features for the predictive conditional AAV% model.

Figure 3 gives the parallel importance view for the conditional AAV% model. Receiving value, target share, prior market price, snap role, and term indicators drive compensation more directly than they drive term.

Table 13: Class-specific sparsity in the lasso multinomial term model.

Term Class	Nonzero Coefficients
1	67
2	45
3	22
4	9
5+	13

Table 13 reports how many transformed predictors remain active for each term class in the lasso multinomial model. The class-specific sparsity pattern is useful because it shows that the interpretable model is not using the same number of signals for every contract length.

Table 14: Ten largest absolute nonzero lasso coefficients by term class.

Term Class	Feature	Coefficient	Absolute Value
1	age At Signing	0.319	0.319
1	prev Contract Start Year	0.281	0.281
1	reg Previous 1 Games	-0.189	0.189
1	prev Contract End Year	0.167	0.167
1	reg Previous 2 Games	-0.166	0.166
1	reg Previous 1 Receiving First Downs	-0.157	0.157
1	prior Team Same yes	-0.156	0.156
1	reg Previous 1 Receiving Yards Per Game	-0.153	0.153
1	ngs Weighted Percentent Share Of Intended Air Yards	-0.142	0.142
1	market Delta Contract Count	-0.128	0.128
2	prior Team Same unknown	0.253	0.253
2	gap Since Prev Contract End	-0.177	0.177
2	reg Weighted Games	-0.162	0.162
2	undrafted Flag yes	0.137	0.137
2	reg Previous 1 Receiving Yards Per Game	-0.122	0.122
2	team Delta Carries	-0.113	0.113
2	reg Previous 1 Receptions Per Game	-0.112	0.112
2	team Previous 1 Games	0.106	0.106
2	reg Weighted Receiving Epa Per Game	-0.104	0.104
2	market Previous 2 Median AAV	-0.094	0.094
3	prev Signed Team unknown	0.662	0.662
3	player Ref Age Gap	0.372	0.372
3	combine Available unknown	0.269	0.269
3	team Weighted Receiving Air Yards	0.161	0.161
3	ngs Previous 1 Avg Yac Above Expectation	0.110	0.110
3	ngs Previous 1 Avg Cushion	0.095	0.095
3	prev Contract AAV	0.037	0.037
3	prev1Team CIN	-0.030	0.030
3	prev1Team NYJ	-0.027	0.027
3	prev1Team LA	0.027	0.027
4	prior Team Same yes	0.183	0.183
4	reg Previous 1 Receiving Epa	0.169	0.169
4	snap Weighted Offense Snaps	0.104	0.104
4	ngs Weighted Yards	0.102	0.102
4	team Previous 1 Receiving Air Yards	-0.078	0.078
4	reg Previous 1 Rushing Yards	0.065	0.065
4	team Weighted Receiving Yards Per Target	0.043	0.043
4	combine Shuttle	0.036	0.036
4	ngs Weighted Yards Per Target	0.004	0.004
5+	start Year	-0.403	0.403
5+	reg Weighted Team Yards Share Calc	0.221	0.221
5+	reg Previous 1 Receiving First Downs Per Game	0.196	0.196
5+	reg Previous 2 Receiving Epa Per Game	0.140	0.140
5+	market Previous 2 Contract Count	-0.085	0.085
5+	reg Previous 1 Receiving Yards	0.030	0.030
5+	reg Previous 1 Air Yards Share	0.019	0.019
5+	team Delta Receiving Yards Per Target	0.018	0.018
5+	team Weighted Losses	-0.012	0.012
5+	team Previous 2 Receiving Air Yards	-0.011	0.011

Table 14 lists the ten largest absolute nonzero lasso coefficients in each class. Coefficients are on the recipe-transformed scale, so they should be read as audit weights rather than raw-unit effects.

Table 15: Largest absolute coefficients in the ridge compensation model.

Feature	Coefficient	Absolute Value
term X5	0.004123	0.004123
term X4	0.002808	0.002808
term X3	0.001850	0.001850
prev Signed Team unknown	-0.001230	0.001230
prev Contract AAV	0.001094	0.001094
prev Contract AAV Percent	0.001079	0.001079
ngs Previous 2 Yards	0.001076	0.001076
ngs Delta Avg Separation	-0.000977	0.000977
ngs Previous 1 Yards	0.000937	0.000937
snap Previous 2 Offense Pct Mean	0.000926	0.000926
prior Team Same yes	0.000908	0.000908
reg Years Available	-0.000882	0.000882
reg Previous 2 Available yes	-0.000868	0.000868
reg Weighted Games	-0.000859	0.000859
reg Previous 1 Receiving Yards	0.000828	0.000828
reg Previous 1 Receiving Yards Per Game	0.000819	0.000819
ngs Previous 1 Avg Yac Above Expectation	0.000814	0.000814
undrafted Flag yes	-0.000814	0.000814
reg Weighted Receiving Tds	0.000811	0.000811
reg Previous 1 Available yes	-0.000808	0.000808
team Years Available	0.000761	0.000761
ngs Previous 1 Receptions	0.000756	0.000756
reg Previous 2 Receiving Tds	0.000746	0.000746
snap Previous 2 Offense Snaps Per Game	0.000734	0.000734
reg Previous 1 Receiving Epa	0.000733	0.000733

Table 15 gives the 25 largest absolute coefficients from the ridge compensation model. Ridge keeps all transformed predictors but shrinks them, so the largest weights are the best compact summary of the linear audit model.

Appendix D: Forecast Outputs

Table 16: Market-level summary of unsigned 2026 WR projections from the predictive pair.

Quantity	Value
Scored unsigned players	163
Players with modal one-year term	163
Mean one-year probability	85.8%
Mean multi-year probability	14.2%
Mean expected AAV	\$1.49M
Median expected AAV	\$1.11M
Maximum expected AAV	\$11.52M

Table 16 summarizes the final scoring run. It is a market-level summary, not a claim that every individual contract will be one year; the probabilities still leave room for player-specific multi-year outcomes.

Table 17: Twenty highest expected-AAV projections from the predictive model pair.

Player	Age	Previous Team	Expected Term	P(1 Year)	P(2+ Years)	Expected AAV	Modal AAV
Deebo Samuel	30	WAS	1.66	63.9%	36.1%	\$11.52M	\$10.16M
Jauan Jennings	28	SF	1.71	62.3%	37.7%	\$10.48M	\$9.05M
Keenan Allen	33	LAC	1.85	57.1%	42.9%	\$10.31M	\$8.73M
Sterling Shepard	33	TB	1.34	77.4%	22.6%	\$4.59M	\$4.19M
DeAndre Hopkins	33	BAL	1.15	90.4%	9.6%	\$4.43M	\$4.19M
Adam Thielen	35	CAR	1.09	93.7%	6.3%	\$4.35M	\$4.22M
Josh Reynolds	31	NYJ	1.05	96.0%	4.0%	\$4.06M	\$4.01M
Ray-Ray McCloud	29	ATL	1.13	90.4%	9.6%	\$3.07M	\$2.90M
Gabriel Davis	27	BUF	1.20	87.8%	12.2%	\$2.69M	\$2.36M
Elijah Moore	25	BUF	1.09	94.3%	5.7%	\$2.67M	\$2.51M
Mike Woods	26	CLE	1.13	89.8%	10.2%	\$2.62M	\$2.52M
Mike Williams	31	LAC	1.12	91.5%	8.5%	\$2.54M	\$2.44M
D'Wayne Eskridge	28	MIA	1.20	86.7%	13.3%	\$2.46M	\$2.35M
Tim Patrick	32	DET	1.14	89.4%	10.6%	\$2.34M	\$2.15M
Erik Ezukanma	26	MIA	1.18	85.0%	15.0%	\$2.14M	\$2.07M
Bo Melton	26	GB	1.33	72.3%	27.7%	\$2.11M	\$2.00M
Marquez Callaway	27	LV	1.09	92.7%	7.3%	\$2.10M	\$2.07M
Zay Jones	30	ARI	1.06	95.7%	4.3%	\$1.96M	\$1.92M
Scotty Miller	28	PIT	1.27	82.8%	17.2%	\$1.93M	\$1.81M
Kristian Wilkerson	29	BUF	1.05	95.9%	4.1%	\$1.90M	\$1.88M

Table 17 expands the main-text forecast discussion to the 20 highest expected-AAV players. Expected AAV is probability-weighted over the term distribution; modal AAV is the AAV attached to the most likely term scenario.


**Figure 4:** One-year probability versus expected AAV for unsigned 2026 wide receivers.

Figure 4 shows the full unsigned market board. The dense cluster at high one-year probability and low AAV is the key market shape; the labeled players are high-price exceptions rather than evidence of a broad multi-year market.




Figure 5: Scenario-level AAV paths for the three highest expected-AAV players. Point size is the term probability assigned by the term model.

Figure 5 displays the scenario-price menu for the three highest expected-AAV players. The lines show that the price model raises AAV under longer supplied terms, but point size shows that the term model assigns most probability to the one-year point.

References

- Chen, Tianqi, and Carlos Guestrin. 2016. “XGBoost: A Scalable Tree Boosting System.” *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785–94. <https://arxiv.org/abs/1603.02754>.
- Friedman, Jerome, Trevor Hastie, and Robert Tibshirani. 2010. “Regularization Paths for Generalized Linear Models via Coordinate Descent.” *Journal of Statistical Software* 33 (1): 1–22. <https://www.jstatsoft.org/article/view/v033i01>.
- nflverse. 2026. *Nflreadr: Download Nflverse Data*. <https://nflreadr.nflverse.com/>.
- Spotrac. 2026. *NFL Wide Receiver Contracts*. <https://www.spotrac.com/nfl/contracts/wide-receiver/>.
- tidymodels. 2026. *Tidymodels: Easily Install and Load the Tidymodels Packages*. <https://www.tidymodels.org/>.